

Double-click (or enter) to edit

▼ Image Captioning - DL Assignment2

Assignment Details

1. Group:90

2. Members

- Afrah Khan - 2019AD04096
- Inderdeep Singh - 2019AD04061
- Partha Pratim Saha - 2019AD04100

Import Libraries/Dataset

```
1 import pandas as pd
2 import numpy as np
3 import re
4 import os.path
5 import tensorflow as tf
6 from tensorflow.keras.preprocessing.image import load_img
7 import matplotlib.pyplot as plt
8 import glob
9 from PIL import Image
10 import string
11 import time
12 pd.set_option('display.max_colwidth', None)
```

▼ Checking GPU availability

```
1 print("Version: ", tf.__version__)
2 print("Eager mode: ", tf.executing_eagerly())
3 print("GPU is", "available" if tf.config.list_physical_devices("GPU") else "NOT AVAILAB
4 print("GPU device name:", tf.test.gpu_device_name())
```

```
Version: 2.5.0
Eager mode: True
GPU is available
GPU device name: /device:GPU:0
```

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

▼ Data Visualization and augmentation:

▼ Read the pickle file

```
1 pkl_df = pd.read_pickle("/content/set_4.pkl")

1 pkl_list=[]
2 for d in pkl_df:
3     pkl_list.append(re.split('#|\t',d,maxsplit=2))

1 mdf = pd.DataFrame(pkl_list,columns =['ImageName', 'id','Caption'])

1 mdf = mdf.groupby(['ImageName']).sum()
2 mdf.drop(columns=['id'],inplace=True)
3 mdf.reset_index(inplace=True)

1 mdf.head()
```

	ImageName	Caption
0	1000268201_693b08cb0e.jpg	A girl go into a wooden building .A child in a pink dress be climb up a set of stair in an entry way .A little girl climb into a wooden playhouse .A little girl climb the stair to her playhouse .A little girl in a pink dress go into a wooden cabin .
1	1001773457_577c3a7d70.jpg	A black dog and a tri-colored dog play with each other on a road .A black dog and a white dog with brown spot be stare at each other in a street .Two dog on pavement move toward each other .
2	1002674143_1b742ab4b8.jpg	There be a girl with pigtail sit in front of a rainbow paint .A little girl be sit in front of a large painted rainbow .
		A man lay on a bench while his dog sit by him .A man lay on a

▼ Image dataset

```
1 import os
2 os.getcwd()

'/content/drive/My Drive/Flicker8k_Dataset'
```

Convert the data into the correct format which could be used for ML model.

```
1 from tqdm import tqdm
2 image_dict = {}
3 for filename in tqdm(glob.glob('*.jpg')):
4     img=load_img(filename,target_size=(224,224))
5     image_dict[filename]=img
```

100% | ██████████ | 8091/8091 [46:45<00:00, 2.88it/s]

```
1 len(image_dict)
```

```
8091
```

```
1 dic={}
2 for i in mdf.ImageName:
3     for j in image_dict:
4         if i in j:
5             dic[i]=image_dict[j]
```

```
1 df_test = pd.DataFrame.from_dict(dic,orient='index',columns=['Image'])
2 df_test.reset_index(inplace=True)
3 df_test.rename(columns={'index':'ImageName'},inplace=True)
```

```
1 main_df = pd.merge(mdf,df_test,on=['ImageName'])
```

```
1 main_df.head()
```

	ImageName	Caption	Image
0	1000268201_693b08cb0e.jpg	A girl go into a wooden building .A child in a pink dress be climb up a set of stair in an entry way .A little girl climb into a wooden playhouse .A little girl climb the stair to her playhouse .A little girl in a pink dress go into a wooden cabin .	<PIL.Image.Image image mode=RGB size=224x224 at 0x7F069F0719D0>
1	1001773457_577c3a7d70.jpg	A black dog and a tri-colored dog play with each other on a road .A black dog and a white dog with brown spot be stare at each other in a street .Two dog on pavement move toward each other .	<PIL.Image.Image image mode=RGB size=224x224 at 0x7F069EEF1550>
2	1002674143_1b742ab4b8.jpg	There be a girl with pigtail sit in front of a rainbow paint .A little girl be sit in front of a	<PIL.Image.Image image mode=RGB size=224x224 at

```
1 main_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8032 entries, 0 to 8031
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ImageName    8032 non-null   object
1   Caption      8032 non-null   object
2   Image        8032 non-null   object
dtypes: object(3)
memory usage: 251.0+ KB
```

Plot at least two samples and their captions (use matplotlib/seaborn/any other library)

```
1 for i in range(len(main_df.head())):
2     print("Caption: " + str(main_df['Caption'][i].split(".")[:-1]))
3     plt.imshow(main_df['Image'][i])
4     plt.axis(False)
5     plt.show()
```

Caption: ['A child in a pink dress be climb up a set of stair in an entry way ', 'A ']



Caption: ['Two dog of different breed look at each other on a road ']



▼ Bring the train and test data in the required format

```
1 image = main_df['Image']
2 for i in range(len(image)):
3     image[i]=np.asarray(image[i])
```

```
1 from sklearn.model_selection import train_test_split
2 # keeping 50% size, as colab session crashes if we increase the train size
3 train, test = train_test_split(main_df, test_size=0.5)
```

```
1 caption_dict={}
2 for i in range(len(train)):
3     caption_dict[train.ImageName.values[i]]=["START "+sent.strip()+" END" for sent in tra
```

```
1 caption_dict.get('1000268201_693b08cb0e.jpg')
```

```
['START a girl go into a wooden building END',
 'START a child in a pink dress be climb up a set of stair in an entry way END',
 'START a little girl climb into a wooden playhouse END',
 'START a little girl climb the stair to her playhouse END',
 'START a little girl in a pink dress go into a wooden cabin END']
```

```
1 count_words = {}
2 count=1
3 for k,vv in caption_dict.items():
4     for v in vv:
```

```

5     for word in v.split():
6         if word not in count_words:
7             count_words[word] = count
8             count +=1

```



```
1 len(count_words)
```

```
4453
```



```

1 for k,vv in caption_dict.items():
2     for v in vv:
3         encode=[]
4         for word in v.split():
5             encode.append(count_words[word])
6         caption_dict[k][vv.index(v)]=encode

```

```
1 caption_dict.get('1000268201_693b08cb0e.jpg')
```

```

[[1, 2, 12, 473, 85, 2, 72, 382, 11],
 [1,
  2,
  37,
  4,
  2,
  249,
  119,
  38,
  168,
  158,
  2,
  182,
  17,
  181,
  4,
  223,
  1356,
  457,
  11],
 [1, 2, 331, 12, 168, 85, 2, 72, 1357, 11],
 [1, 2, 331, 12, 168, 9, 181, 78, 210, 1357, 11],
 [1, 2, 331, 12, 4, 2, 249, 119, 473, 85, 2, 72, 1358, 11]]

```

Model Building

▼ Use Pretrained Resnet-50 model trained on ImageNet dataset

```

1 from tensorflow.keras.applications import ResNet50
2 resnet_model = ResNet50()

```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/resnet50/102973440/102967424> [=====] - 3s 0us/step



```
1 from tensorflow.keras.models import Model
```

```
1 image_model = Model(inputs=resnet_model.input, outputs=resnet_model.layers[-2].output)
```

▼ Bring the train in the required format.

```
1 image_feat = {}
2 for im in range(len(train)):
3     img = train['Image'].values[im].reshape(1,224,224,3)
4     pred = image_model.predict(img).reshape(2048,)
5     image_feat[train.ImageName.values[im]] = pred
```

```
1 from tensorflow.keras.utils import to_categorical
2 from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
1 MAX_LEN = 0
2 for k, vv in caption_dict.items():
3     for v in vv:
4         if len(v) > MAX_LEN:
5             MAX_LEN = len(v)
```

```
1 MAX_LEN
```

```
46
```

```
1 VOCAB_SIZE = len(count_words)
2 def generator(photo, caption):
3     n_samples = 0
4     X = []
5     y_in = []
6     y_out = []
7     for k, vv in caption.items():
8         for v in vv:
9             for i in range(1, len(v)):
10                 try:
11                     X.append(photo[k])
12                     in_seq= [v[:i]]
13                     out_seq = v[i]
14                     in_seq = pad_sequences(in_seq, maxlen=MAX_LEN, padding='post', truncati
15                     out_seq = to_categorical([out_seq], num_classes=VOCAB_SIZE+1)[0]
16                     y_in.append(in_seq)
17                     y_out.append(out_seq)
18                 except:
19                     pass
20
21     return X, y_in, y_out
```

```
1 X, y_in, y_out = generator(image_feat, caption_dict)
```

```

1 X = np.asarray(X)
2 y_in = np.asarray(y_in)
3 y_out = np.asarray(y_out)

1 X.shape,y_in.shape,y_out.shape

((143214, 2048), (143214, 46), (143214, 4454))

1 from tensorflow.keras.models import Model, Sequential
2 from tensorflow.keras import layers
3 from tensorflow.keras.layers import Dense
4 from tensorflow.keras.layers import LSTM
5 from tensorflow.keras.layers import Embedding
6 from tensorflow.keras.layers import Dropout
7 from tensorflow.keras.layers import RepeatVector
8 from tensorflow.keras.layers import TimeDistributed
9 from tensorflow.keras.layers import Concatenate

```

▼ Four(4) layered RNN - We are using LSTM as it is an RNN architecture

Add one layer of dropout at the appropriate position and give reasons.

Adding **L2 regularization** to all the **RNN** layers and one layer of 20% dropout - adding before last layer of **LSTM** because they are the one with the greater number of parameters and thus they're likely to excessively co-adapting themselves causing overfitting. However, since it's a stochastic regularization technique, we can really place it everywhere.

Choose the appropriate activation function for all the layers.

Adding **softmax** as appropriate activation function - **Adam** optimizer as it converges really faster and adding **categorical_crossentropy** as loss function with **accuracy** as performance metric

▼ **Give reasons for the choice of learning rate and its value:** Adding learning rate as 0.0001 as it was hypertuned and gave better results

Model Compilation

```

1 embedding_size = 128
2 max_len = MAX_LEN
3 vocab_size = len(count_words) + 1
4 image_model1 = Sequential()
5 image_model1.add(Dense(embedding_size, input_shape=(2048,), activation='relu'))
6 image_model1.add(Dropout(0.2))

```



```

7 image_model1.add(RepeatVector(max_len))
8 language_model = Sequential()
9 language_model.add(Embedding(input_dim=vocab_size, output_dim=embedding_size, input_len
10 language_model.add(LSTM(128, kernel_regularizer=tf.keras.regularizers.l2(l2=0.0001), re
11 language_model.add(Dropout(0.2))
12 language_model.add(TimeDistributed(Dense(embedding_size)))
13 concat = Concatenate()([image_model1.output, language_model.output])2
14 #Add L2 regularization to all the RNN layers.
15 x = LSTM(128, kernel_regularizer=tf.keras.regularizers.l2(l2=0.0001), return_sequences=
16 x = LSTM(128, kernel_regularizer=tf.keras.regularizers.l2(l2=0.0001), return_sequences=
17 out = Dense(vocab_size,activation='softmax')(x)
18 model = Model(inputs=[image_model1.input, language_model.input], outputs = out)
19
20 #Compiling the model with above parameters
21 model.compile(optimizer=tf.keras.optimizers.Adam(1e-3), loss=tf.losses.categorical_cros
22 #Print the model summary
23 model.summary()

```

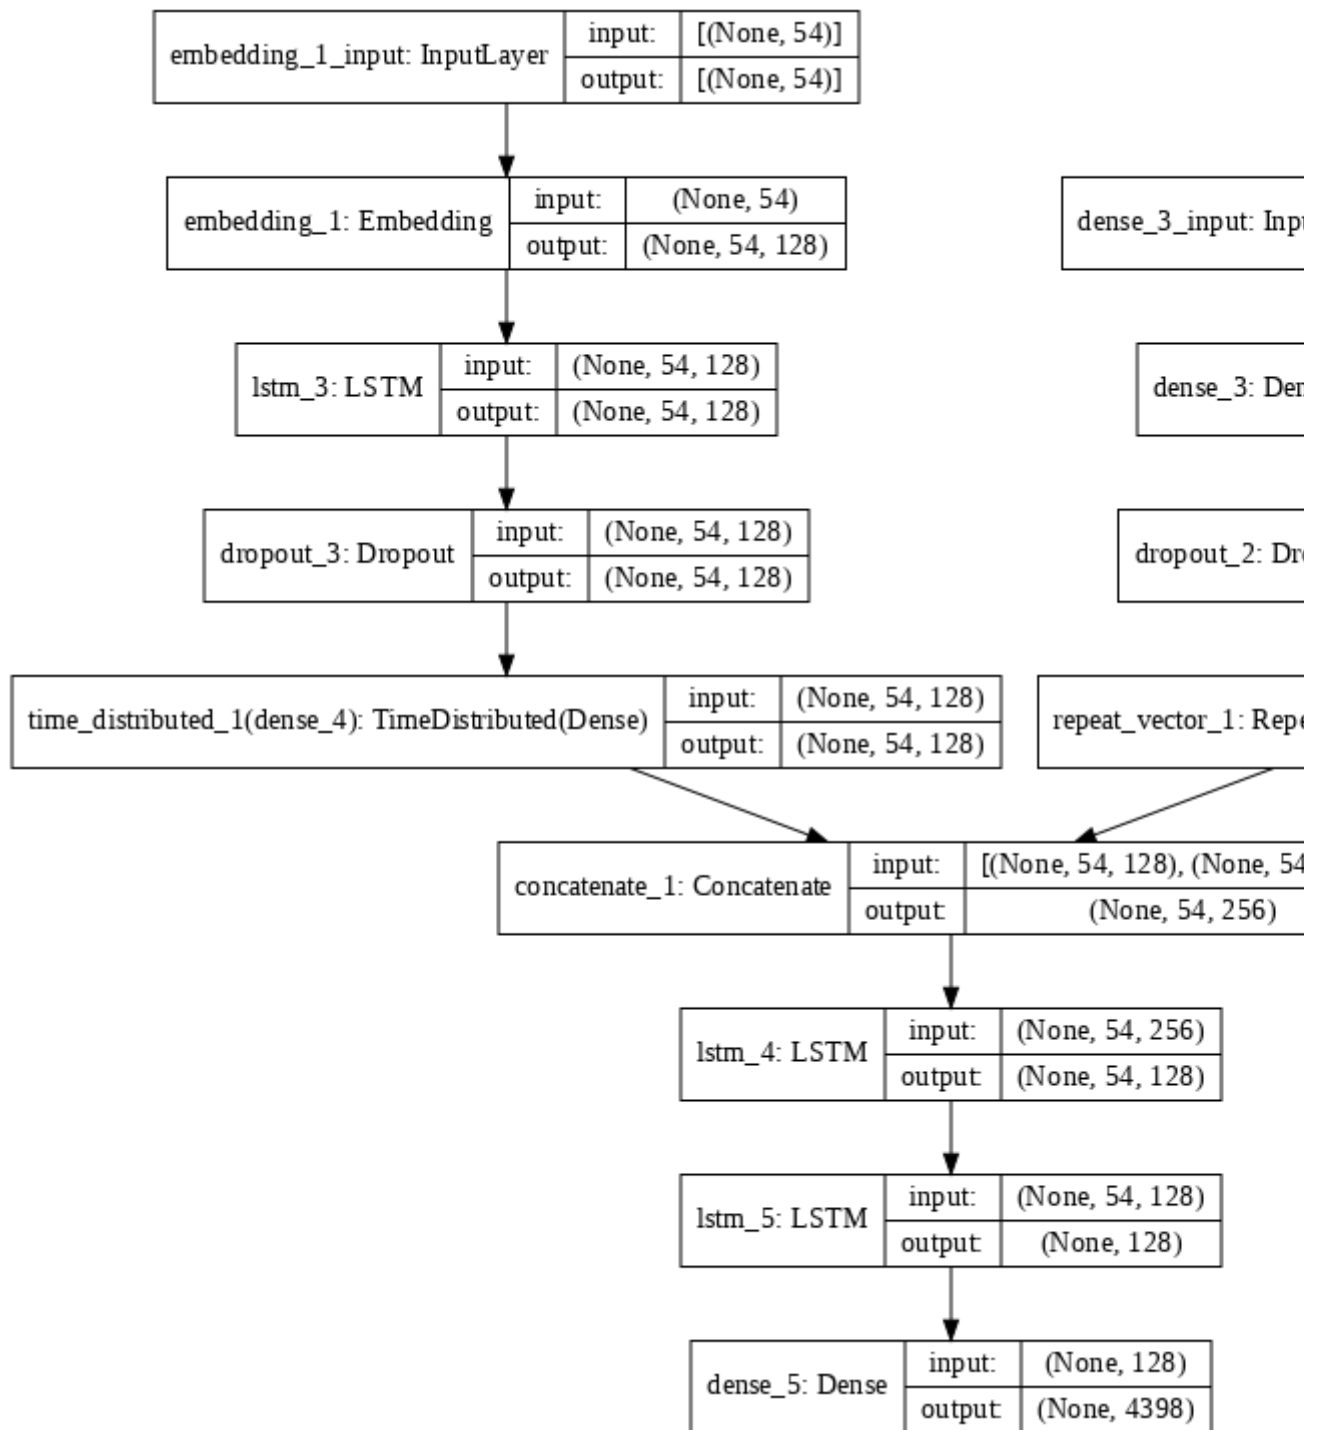
Model: "model_2"

Layer (type)	Output Shape	Param #	Connected to
embedding_1_input (InputLayer)	[(None, 46)]	0	
dense_3_input (InputLayer)	[(None, 2048)]	0	
embedding_1 (Embedding)	(None, 46, 128)	570112	embedding_1_input[0]
dense_3 (Dense)	(None, 128)	262272	dense_3_input[0][0]
lstm_3 (LSTM)	(None, 46, 128)	131584	embedding_1[0][0]
dropout_2 (Dropout)	(None, 128)	0	dense_3[0][0]
dropout_3 (Dropout)	(None, 46, 128)	0	lstm_3[0][0]
repeat_vector_1 (RepeatVector)	(None, 46, 128)	0	dropout_2[0][0]
time_distributed_1 (TimeDistrib	(None, 46, 128)	16512	dropout_3[0][0]
concatenate_1 (Concatenate)	(None, 46, 256)	0	repeat_vector_1[0][0] time_distributed_1[0][0]
lstm_4 (LSTM)	(None, 46, 128)	197120	concatenate_1[0][0]
lstm_5 (LSTM)	(None, 128)	131584	lstm_4[0][0]
dense_5 (Dense)	(None, 4454)	574566	lstm_5[0][0]
Total params: 1,883,750			
Trainable params: 1,883,750			
Non-trainable params: 0			

```

1 from tensorflow.keras.utils import plot_model
2 plot_model(model, show_shapes=True, dpi=72)

```



Model Training

Print the train and validation loss for each epoch. Use the appropriate batch size

```

1 start = time.time()
2 history = model.fit([X, y_in], y_out, batch_size=720, epochs=35, validation_split=0.1,
3 end = time.time()

```

```

Epoch 1/35
180/180 [=====] - 33s 146ms/step - loss: 5.5007 - accuracy: 0.0000
Epoch 2/35
180/180 [=====] - 25s 136ms/step - loss: 5.2176 - accuracy: 0.0000
Epoch 3/35
180/180 [=====] - 25s 137ms/step - loss: 5.1917 - accuracy: 0.0000

```

```

Epoch 4/35
180/180 [=====] - 25s 138ms/step - loss: 5.1286 - accuracy: 0.75
Epoch 5/35
180/180 [=====] - 25s 138ms/step - loss: 5.0974 - accuracy: 0.75
Epoch 6/35
180/180 [=====] - 25s 136ms/step - loss: 5.0780 - accuracy: 0.75
Epoch 7/35
180/180 [=====] - 25s 136ms/step - loss: 5.0485 - accuracy: 0.75
Epoch 8/35
180/180 [=====] - 25s 139ms/step - loss: 5.0293 - accuracy: 0.75
Epoch 9/35
180/180 [=====] - 25s 139ms/step - loss: 5.0088 - accuracy: 0.75
Epoch 10/35
180/180 [=====] - 25s 138ms/step - loss: 4.9890 - accuracy: 0.75
Epoch 11/35
180/180 [=====] - 25s 138ms/step - loss: 4.9679 - accuracy: 0.75
Epoch 12/35
180/180 [=====] - 25s 139ms/step - loss: 4.9458 - accuracy: 0.75
Epoch 13/35
180/180 [=====] - 25s 138ms/step - loss: 4.9252 - accuracy: 0.75
Epoch 14/35
180/180 [=====] - 25s 137ms/step - loss: 4.8974 - accuracy: 0.75
Epoch 15/35
180/180 [=====] - 25s 137ms/step - loss: 4.8752 - accuracy: 0.75
Epoch 16/35
180/180 [=====] - 25s 138ms/step - loss: 4.8574 - accuracy: 0.75
Epoch 17/35
180/180 [=====] - 25s 140ms/step - loss: 4.8331 - accuracy: 0.75
Epoch 18/35
180/180 [=====] - 25s 138ms/step - loss: 4.8143 - accuracy: 0.75
Epoch 19/35
180/180 [=====] - 25s 138ms/step - loss: 4.7929 - accuracy: 0.75
Epoch 20/35
180/180 [=====] - 25s 139ms/step - loss: 4.7757 - accuracy: 0.75
Epoch 21/35
180/180 [=====] - 25s 139ms/step - loss: 4.7597 - accuracy: 0.75
Epoch 22/35
180/180 [=====] - 25s 138ms/step - loss: 4.7465 - accuracy: 0.75
Epoch 23/35
180/180 [=====] - 25s 137ms/step - loss: 4.7344 - accuracy: 0.75
Epoch 24/35
180/180 [=====] - 25s 138ms/step - loss: 4.7160 - accuracy: 0.75
Epoch 25/35
180/180 [=====] - 25s 139ms/step - loss: 4.7021 - accuracy: 0.75
Epoch 26/35
180/180 [=====] - 25s 138ms/step - loss: 4.6927 - accuracy: 0.75
Epoch 27/35
180/180 [=====] - 25s 138ms/step - loss: 4.6763 - accuracy: 0.75
Epoch 28/35
180/180 [=====] - 25s 139ms/step - loss: 4.6613 - accuracy: 0.75
Epoch 29/35
180/180 [=====] - 25s 139ms/step - loss: 4.6546 - accuracy: 0.75

```

total time taken for training: ~15 mins

```

1 # Print the total time taken for training
2 print("Total time taken for training: ", (end - start)/60, "mins")

```


Total time taken for training: 15.453375518321991 mins

Plot the loss and accuracy history graphs for both train and validation set

```

1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5 epochs_range = range(len(history.epoch))
6
7 plt.figure(figsize=(8, 8))
8 plt.subplot(1, 2, 1)
9 plt.plot(epochs_range, acc, label='Training Accuracy')
10 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
11 plt.legend(loc='lower right')
12 plt.title('Training and Validation Accuracy')
13
14 plt.subplot(1, 2, 2)
15 plt.plot(epochs_range, loss, label='Training Loss')
16 plt.plot(epochs_range, val_loss, label='Validation Loss')
17 plt.legend(loc='upper right')
18 plt.title('Training and Validation Loss')
19 plt.show()

```



```
1 inv_dict = {v:k for k, v in count_words.items()}
```

```
1 def getTimage(v):
```

```

1 def getImage(x):
2     test_img = np.reshape(x, (1,224,224,3))
3     return test_img

```

Model Evaluation: Take a random image from google and generate caption for that image

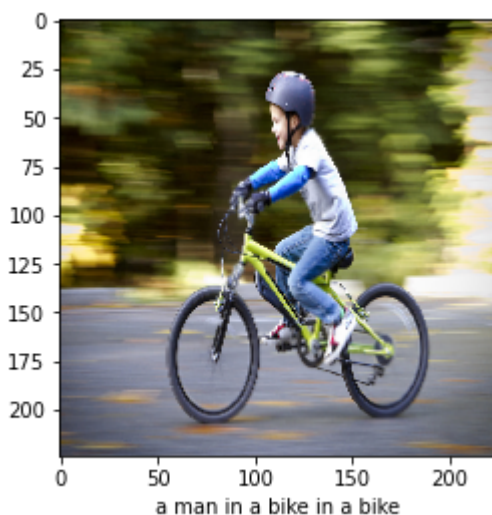
```

1 testImg = load_img('/content/test_image3.jpg',target_size=(224,224,3))

1 test_pred = image_model.predict(getImage(testImg)).reshape(1,2048)
2 text_inp = ['START']
3 count = 0
4 caption = ''
5 while count < 25:
6     count += 1
7     encoded = []
8     for i in text_inp:
9         encoded.append(count_words[i])
10    encoded = [encoded]
11    encoded = pad_sequences(encoded, padding='post', truncating='post', maxlen=
12    prediction = np.argmax(model.predict([test_pred, encoded]))
13    sampled_word = inv_dict[prediction]
14    caption = caption + ' ' + sampled_word
15    if sampled_word == 'END':
16        break
17    text_inp.append(sampled_word)
18 plt.figure()
19 plt.imshow(testImg)
20 plt.xlabel(caption.replace("END","").strip())

```

Text(0.5, 0, 'a man in a bike in a bike')



✓ 1s completed at 23:02 ● ✕